

# An optimization-inspired approach to parallel sorting

Team Metropolis:  
James Farzi, JJ Lay, Graham West

February 13, 2019

## 1 The Algorithm

- Distributing/Importing Files
- Sorting
- Binning
- Exchanging data

## 2 Testing

- Methodology
- Results

## 3 Conclusions

# The Algorithm

# The Algorithm

XXX

- XXX
- XXX

XXX

- XXX
- XXX

# Distributing/Importing Files

XXX

- XXX
- XXX

XXX

- XXX
- XXX

We implemented Merge and Bubble Sort

XXX

- XXX
- XXX

XXX

- XXX
- XXX

## Binary search

- Since the data is sorted, we can use a binary search to find where the bin edges lie in index space
- We can then subtract successive edges' indices to find the number of elements in that bin

## XXX

- XXX
- XXX

## Adapting the bins

for interior bin edges (endpoint bins stay constant):

$$\begin{aligned}\Delta C &= 2.0(c_i^n - c_{i-1}^n)/(c_i^n + c_{i-1}^n) \\ \Delta B &= b_{i+1}^n - b_i^n \\ b_i^{n+1} &= b_i^n + \alpha \Delta C \Delta B\end{aligned}\tag{1}$$

where  $0 < \alpha < 0.5$  and  $b_i^n < b_{i+1}^n$  for all  $n$

## Uniformity metric

$$U^n = \max\left(\frac{c_{\max} - c_{\text{avg}}}{c_{\text{avg}}}, \frac{c_{\text{avg}} - c_{\min}}{c_{\text{avg}}}\right)\tag{2}$$



# Binning

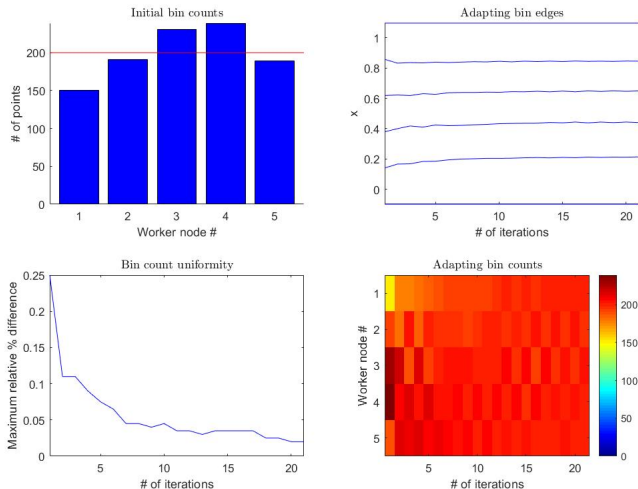


Figure 1: Using 5 nodes to uniformly bin 1000 data points

# Exchanging data

XXX

- XXX
- XXX

XXX

- XXX
- XXX

# Testing

# Conclusions